

LAMBTON COLLEGE



A Report on [Lab 4,5,6 on AWS Cloud Foundations]

121 Brunel Rd, Mississauga

ON L4Z 3E9

A Group assignment with screenshots of Lab 4, 5, and 6

on Aws academy

Big Data Analytics DSMM

**Under the supervision
of
Professor Pedram Habibi**

Submitted BY:

Aadarsha Chapagain (C0825975)

Davinderjit Singh (C0833117)

Priti Bhale (C0835691)

Milanjeet Kaur(C0829899)

Palwinder Kaur (C0827804)

Nimmo Usman (C0836309)

Submitted To:

Lambton College

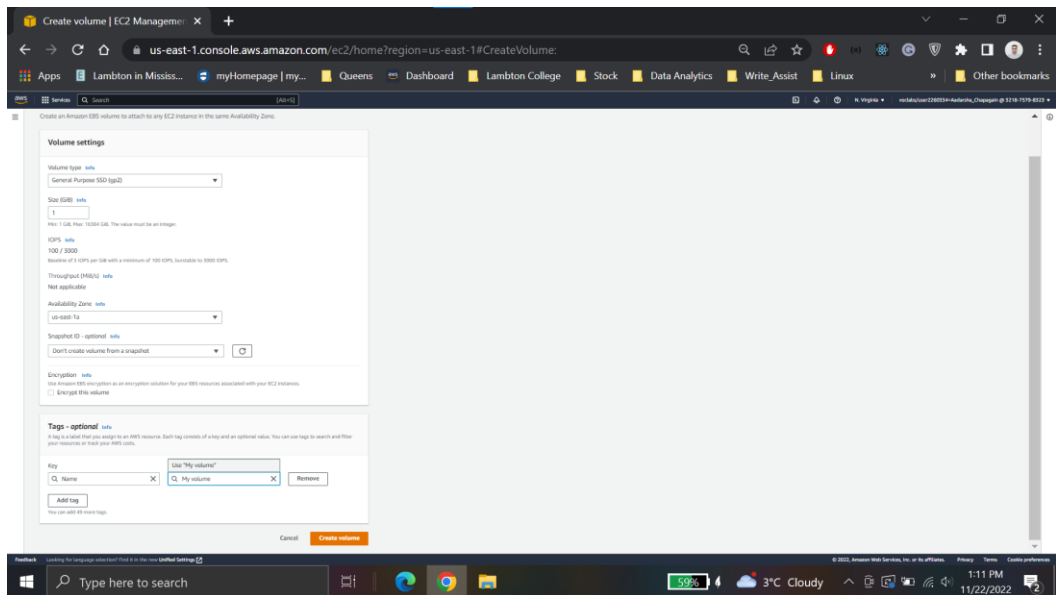
Professor Pedram Habibi

Submission Date:

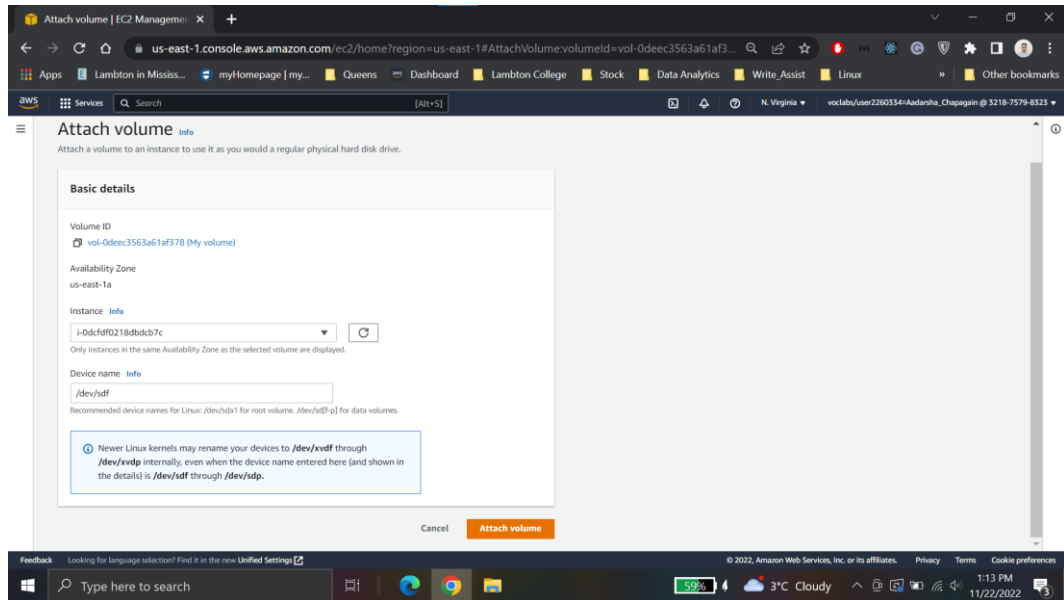
22nd November 2022

Lab4: Working With EBS

Create Volume



Attach Volume



Login using Putty client

At first the .ppk was downloaded and using same key in putty ssh connection was established.

[illegible]

Create an ext3 file system on the new volume

```
ec2-user@ip-10-1-11-190:~$
# login as: ec2-user
# Authenticating with public key "imported-openssh-key"

      _   _
     / \   \
    /___\   \
   /___/ \___\
  /___/   \___\

Amazon Linux 2 AMI

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-10-1-11-190 ~]$ df -h
filesystem      Size  Used Avail Use% Mounted on
devtmpfs        484M   0  484M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M 412K  492M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G 1.5G  6.6G 19% /
tmpfs           99M   0   99M   0% /run/user/1000
[ec2-user@ip-10-1-11-190 ~]$ sudo mkfs -t ext3 /dev/sdf
mkfs2fs 1.42.9 (28-Dec-2013)
Filesystem label=
OS type: Linux
Block size=4096 (log=2)
Fragment size=4096 (log=2)
Stride=0 blocks, Stripe width=0 blocks
8536 inodes, 262144 blocks
13107 blocks (5.00%) reserved for the super user
First data block=0
Maximum filesystem blocks=268435456
4 block groups
32768 blocks per group, 32768 fragments per group
8192 inodes per group
Superblock backups stored on blocks:
        32768, 96304, 163840, 229376

Allocating group tables: done
Writing inode tables: done
creating journal (8192 blocks): done
Writing superblocks and filesystem accounting information: done

[ec2-user@ip-10-1-11-190 ~]$
```

Create a directory for mounting the new storage volume:

```
ec2-user@ip-10-1-11-190:~$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        484M   0  484M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M  412K  491M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.5G   6.6G  19% /
tmpfs           99M   0   99M   0% /run/user/1000
[ec2-user@ip-10-1-11-190 ~]$ sudo mkfs -t ext3 /dev/sdf
mke2fs 1.42.9 (28-Dec-2013)
Filesystem label=
OS type: Linux
Block size=4096 (log=2)
Fragment size=4096 (log=2)
Stride=0 blocks, Stripe width=0 blocks
65536 inodes, 262144 blocks
13107 blocks (5.00%) reserved for the super user
First data block=0
Maximum filesystem blocks=268435456
8 block groups
32768 blocks per group, 32768 fragments per group
8192 inodes per group
Superblock backups stored on blocks:
32768, 98304, 163840, 229376

Allocating group tables: done
Writing inode tables: done
Creating journal (8192 blocks): done
Writing superblocks and filesystem accounting information: done

[ec2-user@ip-10-1-11-190 ~]$ dh -h
-bash: dh: command not found
[ec2-user@ip-10-1-11-190 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        484M   0  484M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M  412K  491M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.5G   6.6G  19% /
tmpfs           99M   0   99M   0% /run/user/1000
[ec2-user@ip-10-1-11-190 ~]$ sudo mkdir /mnt/data-store
[ec2-user@ip-10-1-11-190 ~]$
```

Mount the new volume:

```
ec2-user@ip-10-1-11-190:~$ sudo mkdir /mnt/data-store
ec2-user@ip-10-1-11-190:~$ sudo mount /dev/sdf /mnt/data-store
ec2-user@ip-10-1-11-190:~$
```

To configure the Linux instance to mount this volume whenever the instance is started, add a line to */etc/fstab*.

```
[ec2-user@ip-10-1-11-190 ~]$ echo "/dev/sdf /mnt/data-store ext3 defaults,noatime 1 2" | sudo tee -a /etc/fstab
/dev/sdf /mnt/data-store ext3 defaults,noatime 1 2
[ec2-user@ip-10-1-11-190 ~]$
```

View the configuration file to see the setting on the last line

```
[ec2-user@ip-10-1-11-190 ~]$ cat /etc/fstab
#
UUID=f71e0200-1724-4492-b770-0f9d19cb102a / xfs defaults,noatime 1 1
/dev/sdf /mnt/data-store ext3 defaults,noatime 1 2
[ec2-user@ip-10-1-11-190 ~]$
```

View the available storage again

```
[ec2-user@ip-10-1-11-190 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        484M   0  484M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M  412K  491M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.5G   6.6G  19% /
tmpfs           99M   0   99M   0% /run/user/1000
/dev/xvdf       975M   60K  924M   1% /mnt/data-store
[ec2-user@ip-10-1-11-190 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        484M   0  484M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M  412K  491M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.5G   6.6G  19% /
tmpfs           99M   0   99M   0% /run/user/1000
/dev/xvdf       975M   60K  924M   1% /mnt/data-store
[ec2-user@ip-10-1-11-190 ~]$
```

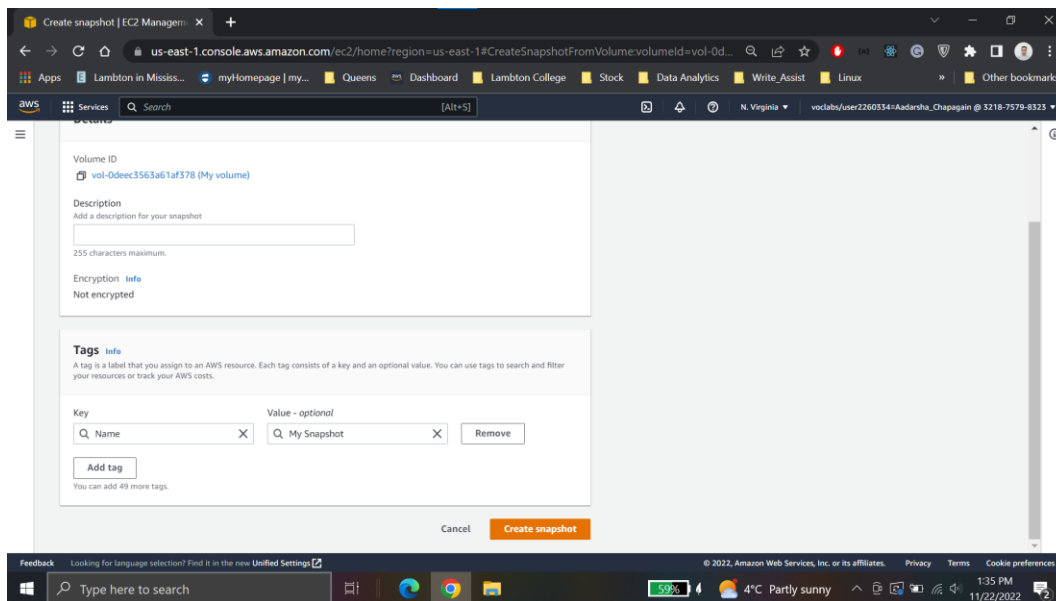
On mounted volume, create a file and add some text to it.

```
/dev/xvdi          512M  BkB  924M  i/ /mnt/data-store
[ec2-user@ip-10-1-11-190 ~]$ sudo sh -c "echo some text has been written > /mnt/data-store/file.txt"
>
[ec2-user@ip-10-1-11-190 ~]$
```

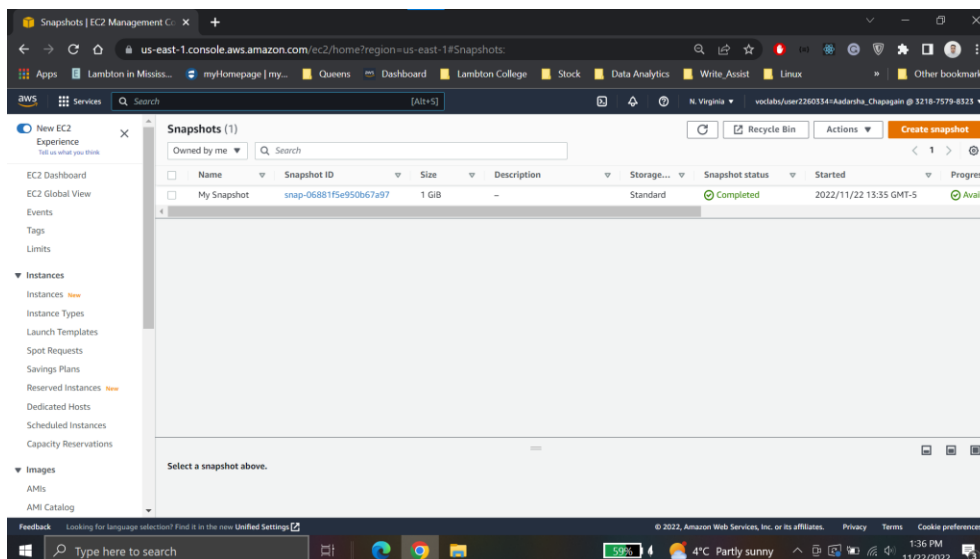
Verify that the text has been written to your volume

```
>
[ec2-user@ip-10-1-11-190 ~]$ cat /mnt/data-store/file.txt
some text has been written
[ec2-user@ip-10-1-11-190 ~]$
```

Create Snapshot

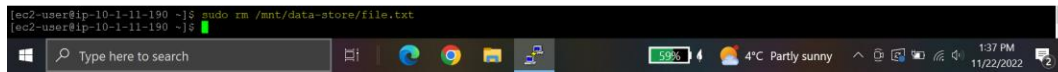


Snapshot created

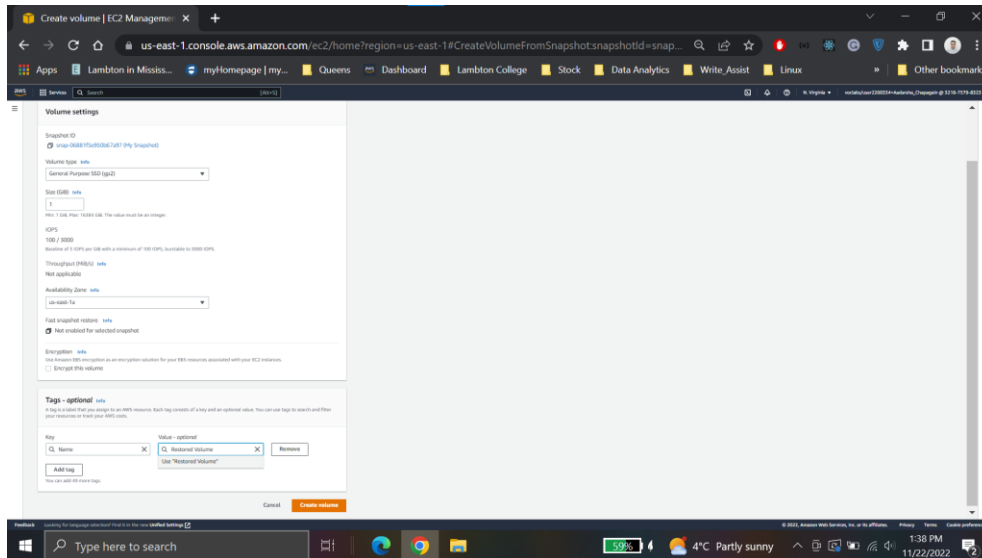


Delete created file

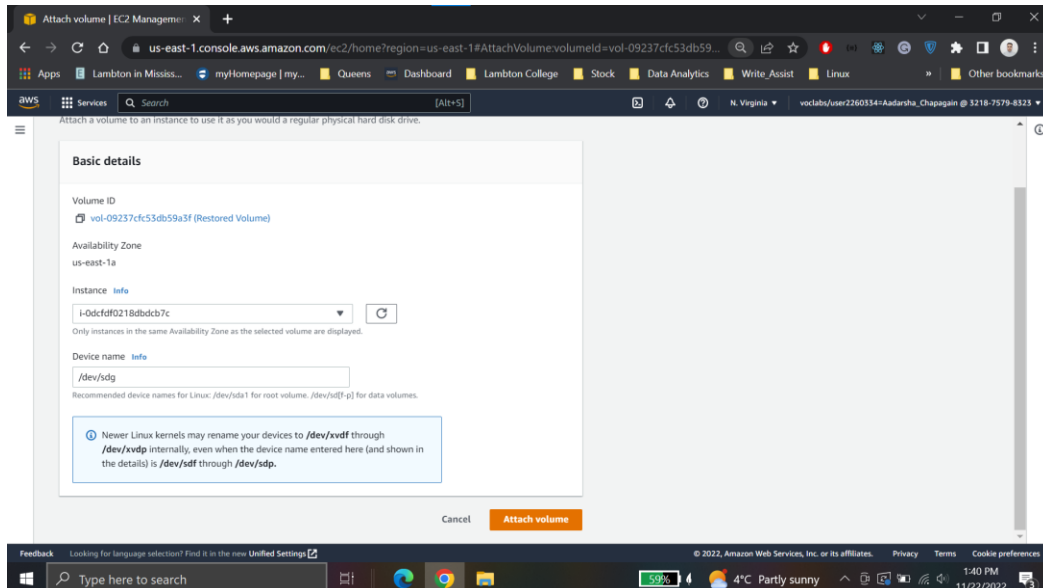
```
ec2-user@ip-10-1-11-190 ~]$ sudo rm /mnt/data-store/file.txt
ec2-user@ip-10-1-11-190 ~]$
```

A terminal window on a Linux EC2 instance. The user runs the command 'sudo rm /mnt/data-store/file.txt' and the prompt returns to the user. The taskbar at the bottom shows the date as 11/22/2022 and the time as 1:37 PM.

Create volume from stored snapshot



Attach the volume to instance



Create a directory for mounting the new storage volume:

sudo mkdir /mnt/data-store2

Mount the new volume:

sudo mount /dev/sdg /mnt/data-store2

```
ec2-user@ip-10-1-11-190 ~]$ cat /mnt/data-store/file.txt
some text has been written
ec2-user@ip-10-1-11-190 ~]$ sudo rm /mnt/data-store/file.txt
ec2-user@ip-10-1-11-190 ~]$ ls /mnt/data-store/
lost+found
ec2-user@ip-10-1-11-190 ~]$ sudo mkdir /mnt/data-store2
ec2-user@ip-10-1-11-190 ~]$ sudo mount /dev/sdg /mnt/data-store2
ec2-user@ip-10-1-11-190 ~]$ ls /mnt/data-store2/
ec2-user@ip-10-1-11-190 ~]$ sudo mount /dev/sdg /mnt/data-store2
ec2-user@ip-10-1-11-190 ~]$ ls /mnt/data-store2/
file.txt lost+found
ec2-user@ip-10-1-11-190 ~]$
```

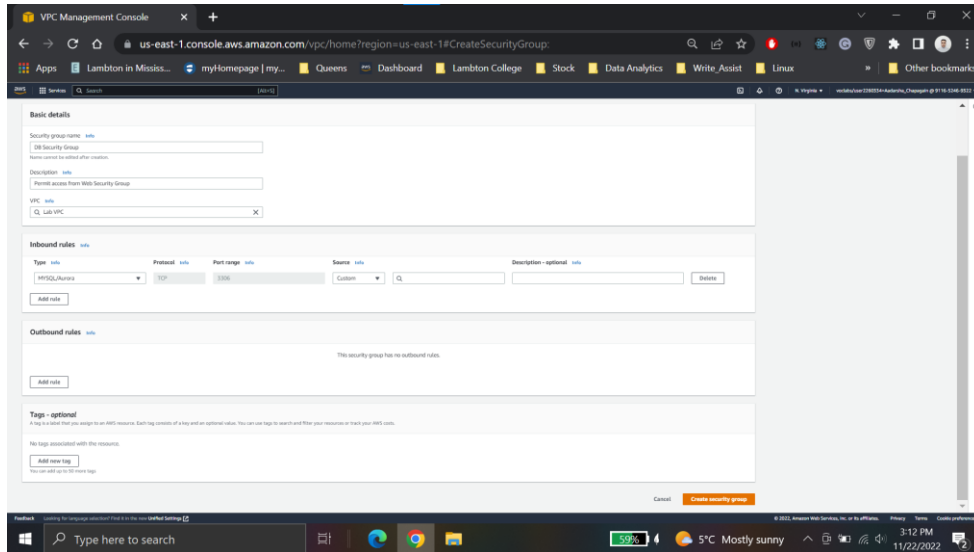


Conclusions for Lab4

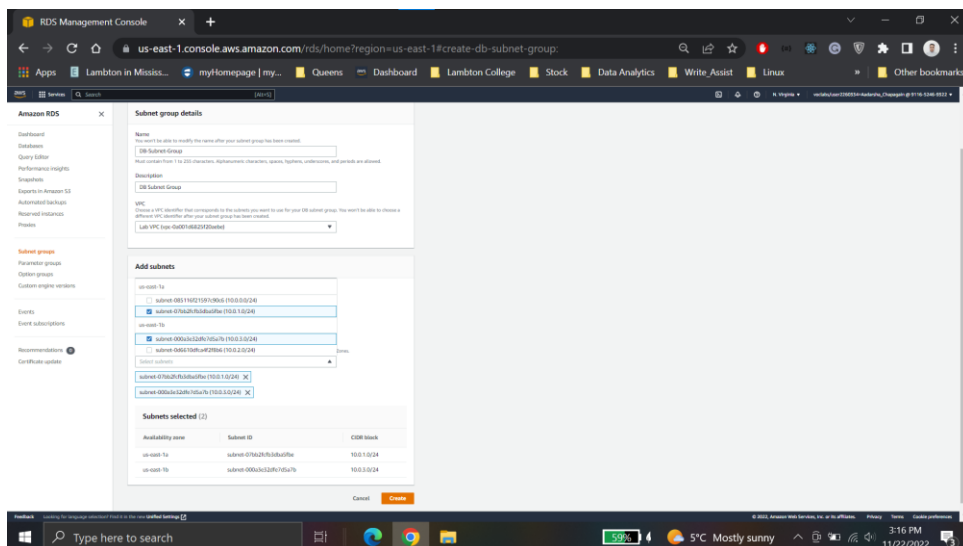
- Created an Amazon EBS volume
- Attached the volume to an EC2 instance
- Created a file system on the volume
- Added a file to volume
- Created a snapshot of your volume
- Created a new volume from the snapshot
- Attached and mounted the new volume to your EC2 instance
- Verified that the file created earlier was on the newly created volume

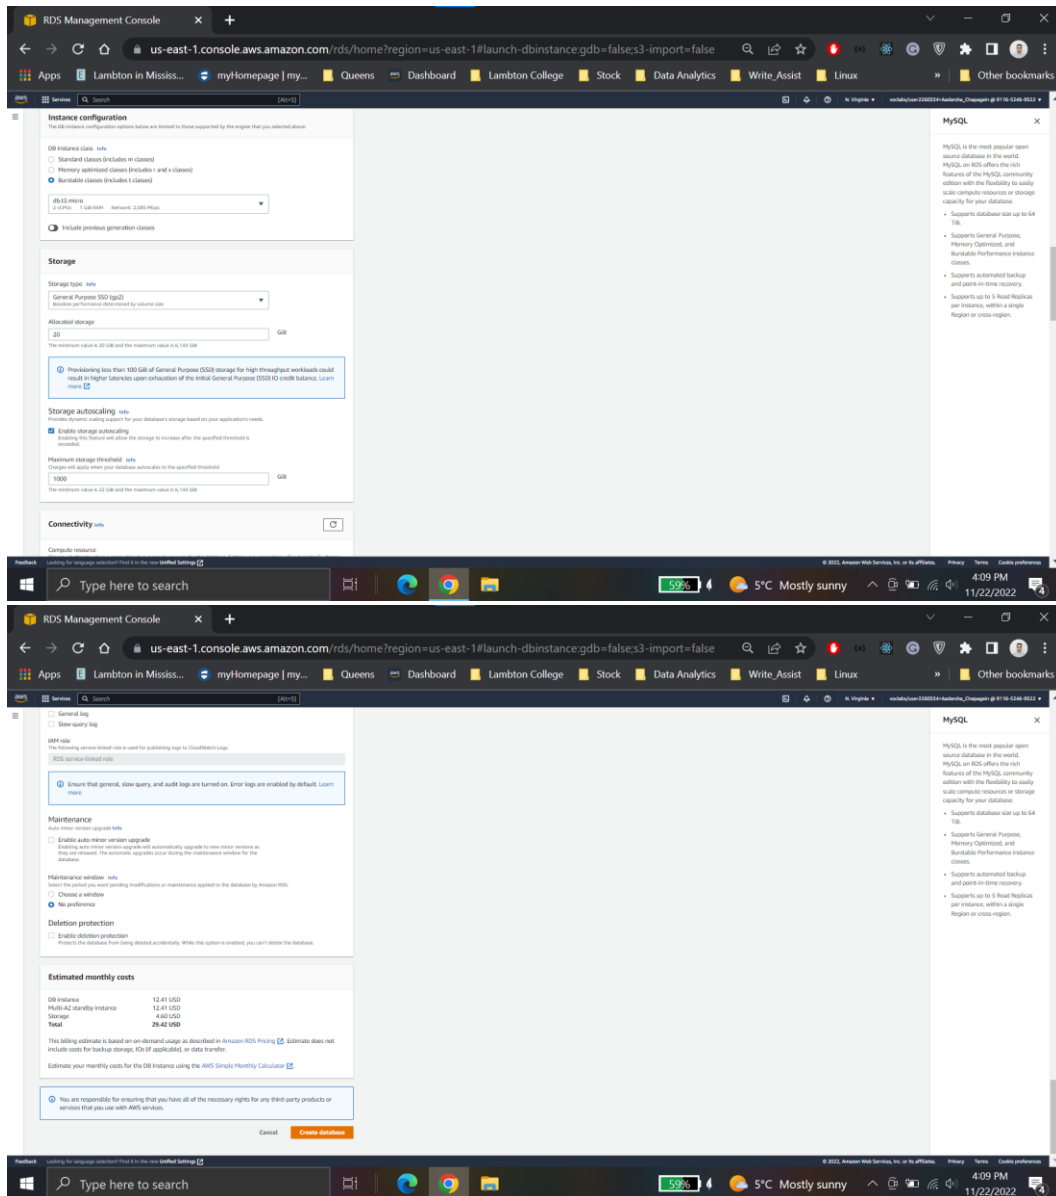
Lab 5: Build Your DB Server and Interact With Your DB Using an App

Create a security group




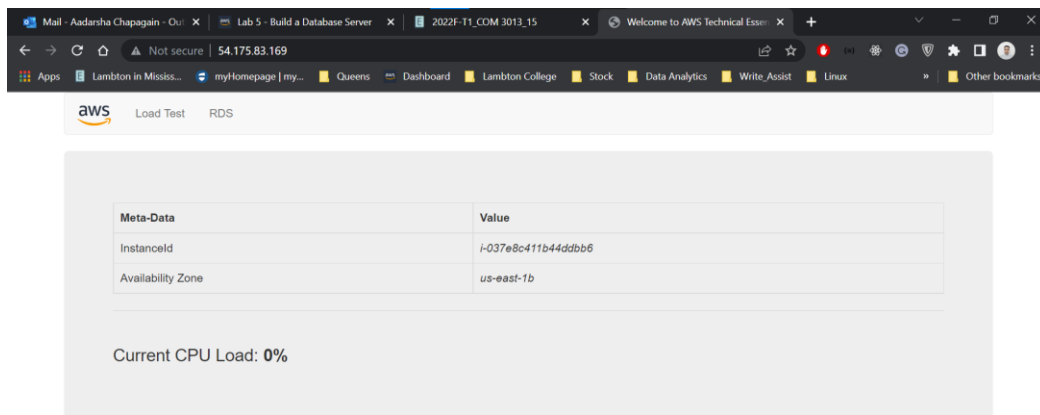
Create Db Subnet Group



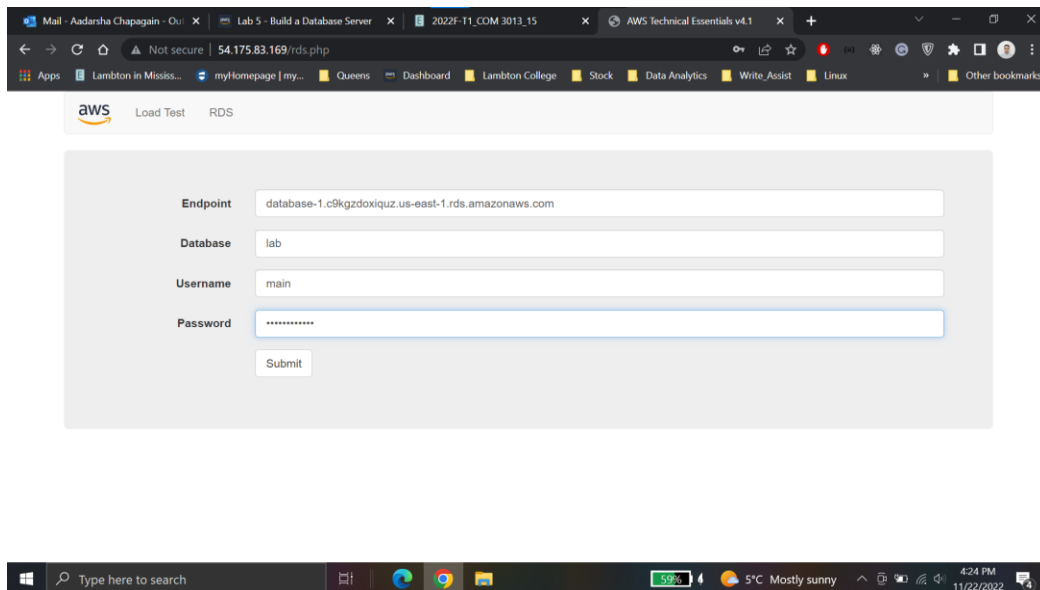


database-1.c9kgzdxiqz.us-east-1.rds.amazonaws.com

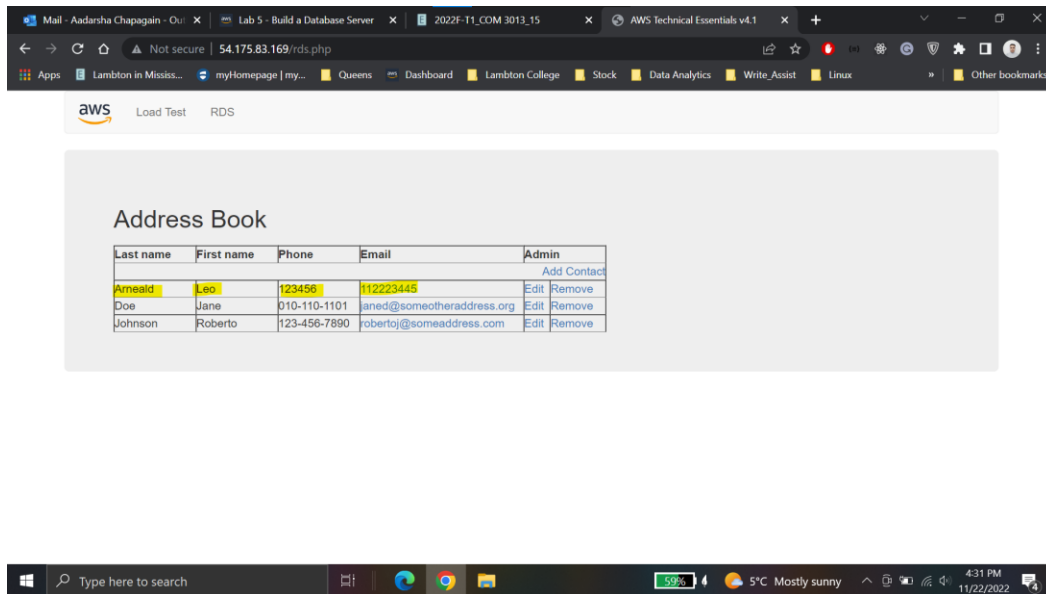
Web server Ip address



Web application to interact with database.



Created the record to test the connection



Lab 5 Conclusion

- Launched an Amazon RDS DB instance with high availability.
- Configured the DB instance to permit connections from your web server.
- Opened a web application and interact with your database

Lab 6: Scale and Load Balance Your Architecture

Task 1: Create an AMI for Auto Scaling

The screenshot shows the 'Create image' page in the AWS Management Console. The page title is 'Create image' with an 'Info' link. Below the title, it states: 'An image (also referred to as an AMI) defines the programs and settings that are applied when you launch an EC2 instance. You can create an image from the configuration of an existing instance.' The 'Instance ID' field is populated with 'i-088763f4dc8dae34a (Web Server 1)'. The 'Image name' field contains 'WebServerAMI'. The 'Image description - optional' field contains 'Lab AMI for Web Server'. There is an unchecked checkbox for 'No reboot' and an unchecked checkbox for 'Enable'. Below these are 'Instance volumes' and a table with columns: Volume type, Device, Snapshot, Size, Volume type, IOPS, Throughput, Delete on, and Encry. The bottom of the console shows the AWS logo, a search bar, and a navigation menu. The system tray at the bottom indicates the date is 11/22/2022 and the time is 6:07 PM.

Create image [Info](#)

An image (also referred to as an AMI) defines the programs and settings that are applied when you launch an EC2 instance. You can create an image from the configuration of an existing instance.

Instance ID
i-088763f4dc8dae34a (Web Server 1)

Image name
WebServerAMI
Maximum 127 characters. Can't be modified after creation.

Image description - optional
Lab AMI for Web Server
Maximum 255 characters

No reboot
☐ Enable

Instance volumes

Volume type	Device	Snapshot	Size	Volume type	IOPS	Throughput	Delete on	Encry
-------------	--------	----------	------	-------------	------	------------	-----------	-------

Feedback Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Type here to search

59% 3°C Mostly clear 6:07 PM 11/22/2022

Task 2: Create a Load Balancer

Create a target group

The screenshot shows the 'Create target group' page in the AWS Management Console. The page title is 'Target groups | EC2 Management Console'. The breadcrumb navigation is 'EC2 > Target groups > Create target group'. The page has two tabs: 'Specify group details' and 'Register targets'. The 'Register targets' tab is active. It shows a table of 'Available instances (2)'. The table has columns: Instance ID, Name, State, Security groups, Zone, and Subnet ID. Two instances are listed: 'i-088763f4dc8dae34a' (Web Server 1) and 'i-088763f4dc8dae34a' (Web Server 1). Below the table, there is a section for 'Review targets' with a 'Targets (0)' list and a 'No instances added yet' message. The bottom of the console shows the AWS logo, a search bar, and a navigation menu. The system tray at the bottom indicates the date is 11/22/2022 and the time is 6:16 PM.

Target groups | EC2 Management Console

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#CreateTargetGroup

EC2 > Target groups > Create target group

Specify group details

Register targets

Register targets

This is an optional step to create a target group. However, to ensure that your load balancer routes traffic to this target group you must register your targets.

Available instances (2)

Instance ID	Name	State	Security groups	Zone	Subnet ID
i-088763f4dc8dae34a	Web Server 1	running	sg-088763f4dc8dae34a	us-east-1a	subnet-088763f4dc8dae34a
i-088763f4dc8dae34a	Web Server 1	running	sg-088763f4dc8dae34a	us-east-1a	subnet-088763f4dc8dae34a

0 selected

Ports for the selected instances
Ports for routing traffic to the selected instances
ID
1-65535 (specify multiple ports with commas)
Include as pending below

Review targets

Targets (0)

All Filter instances by property or value

Remove all pending

Remove Health status Instance ID Name Port State Security groups Zone Subnet ID

No instances added yet

Specify instances above, or leave this group empty if you prefer to add targets later.

0 pending

Cancel Previous Create target group

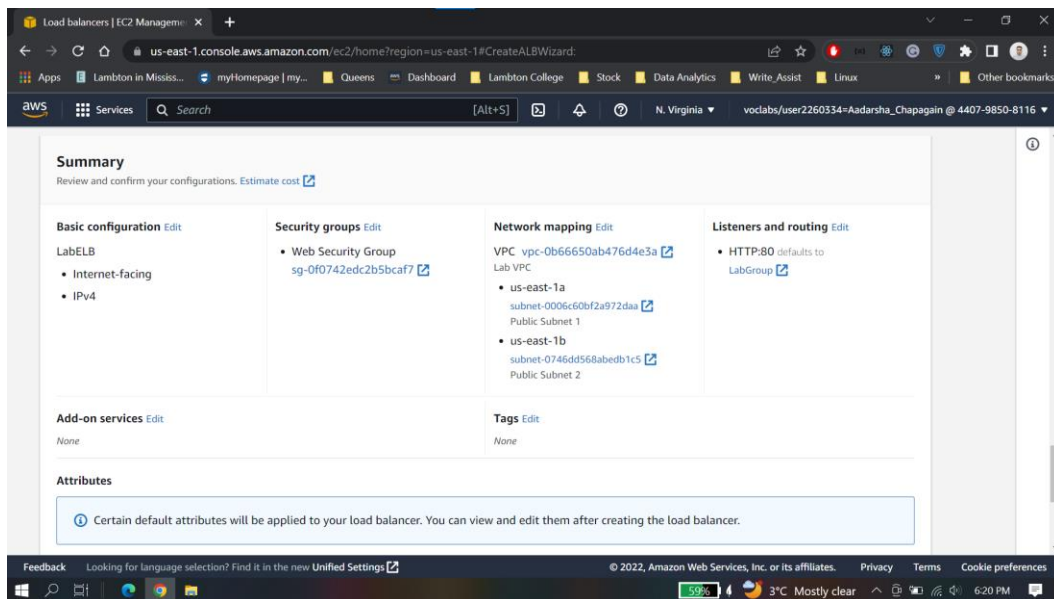
Feedback Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

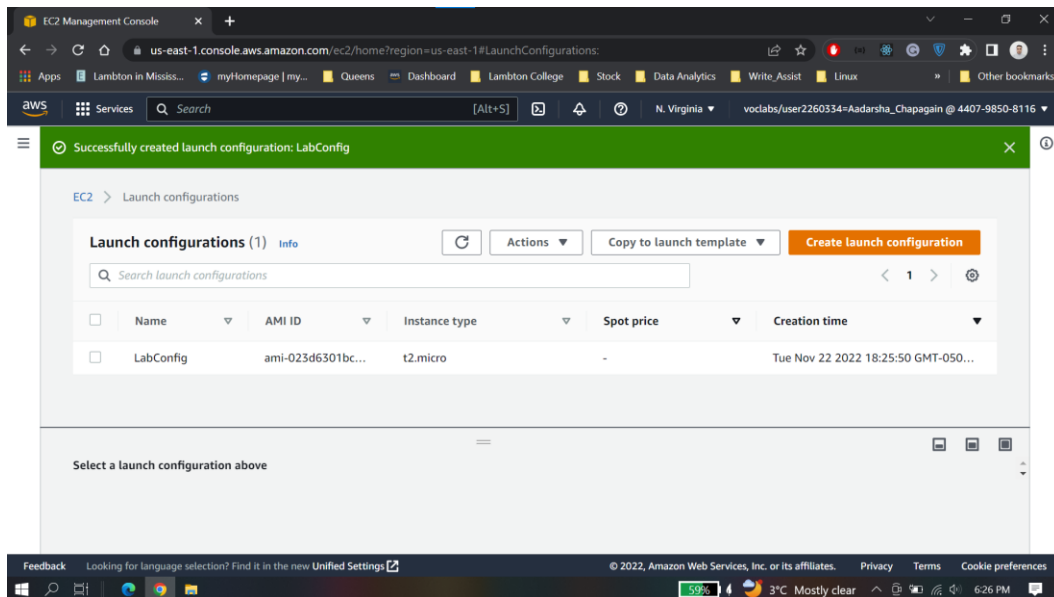
Type here to search

59% 3°C Mostly clear 6:16 PM 11/22/2022

Create Load Balancer



Task 3: Create a Launch Configuration and an Auto Scaling Group



Creta Auto scaling group

Create Auto Scaling group | EC2

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1&redirectFrom=asg#CreateAutoScalingGroup.la...

Apps | Lambda in Mississ... | myHomepage | my... | Queens | Dashboard | Lambton College | Stock | Data Analytics | Write_Assist | Linux | Other bookmarks

Services | Search | [Alt+S]

N. Virginia | voclabs/user2260334-kadasha_Chappan @ 4407-9850-8116

Step 1
Choose launch template or configuration

Step 2
Choose instance launch options

Step 3 (optional)
Configure advanced options

Step 4 (optional)
Configure group size and scaling policies

Step 5 (optional)
Add notifications

Step 6 (optional)
Add tags

Step 7
Review

Configure advanced options

Choose a load balancer to distribute incoming traffic for your application across instances to make it more reliable and easily scalable. You can also set options that give you more control over health check replacements and monitoring.

Load balancing - optional

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

☐ No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.

☒ Attach to an existing load balancer
Choose from your existing load balancers.

☐ Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.

Attach to an existing load balancer

Select the load balancers that you want to attach to your Auto Scaling group.

☒ Choose from your load balancer target groups
This option allows you to attach Application, Network, or Gateway Load Balancers.

☐ Choose from Classic Load Balancers

Existing load balancer target groups

Only instance target groups that belong to the same VPC as your Auto Scaling group are available for selection.

Select target groups:

LabGroup | HTTP
Application Load Balancer: LAB0ELB

Health checks - optional

Step 1
Choose launch template or configuration

Step 2
Choose instance launch options

Step 3 (optional)
Configure advanced options

Step 4 (optional)
Configure group size and scaling policies

Step 5 (optional)
Add notifications

Step 6 (optional)
Add tags

Configure group size and scaling policies

Set the desired, minimum, and maximum capacity of your Auto Scaling group. You can optionally add a scaling policy to dynamically scale the number of instances in the group.

Group size - optional

Specify the size of the Auto Scaling group by changing the desired capacity. You can also specify minimum and maximum capacity limits. Your desired capacity must be within the limit range.

Desired capacity

Minimum capacity

Maximum capacity

Scaling policies - *optional*

Choose whether to use a scaling policy to dynamically resize your Auto Scaling group to meet changes in demand. [Info](#)

- ☒ **Target tracking scaling policy**
Choose a desired outcome and leave it to the scaling policy to add and remove capacity as needed to achieve that outcome.

☐ None

Scaling policy name

LabScalingPolicy

Metric type

Average CPU utilization ▼

Target value

60

Instances need

300

seconds warm up before including in metric

☐ Disable scale in to create only a scale-out policy

Instance scale-in protection - *optional*

Instance scale-in protection

If protect from scale in is enabled, newly launched instances will be protected from scale in by default.

☐ Enable instance scale-in protection

Cancel

Previous

Skip to review

Next

Review the configurations

Review

Info

Step 1: Choose launch template or configuration

Edit

Group details

Auto Scaling group name

Lab Auto Scaling Group

Launch configuration

LabConfig

Step 2: Choose instance launch options

Edit

Network

Network

VPC

vpc-0b66650ab476d4e3a

Availability Zone

Subnet

us-east-1a	subnet-0bc9b7af0d737722b	10.0.1.0/24
us-east-1b	subnet-0cd0249ddcf2c2462	10.0.3.0/24

Newly launched two instances

Instances (4)

Info

Connect

Instance state

Actions

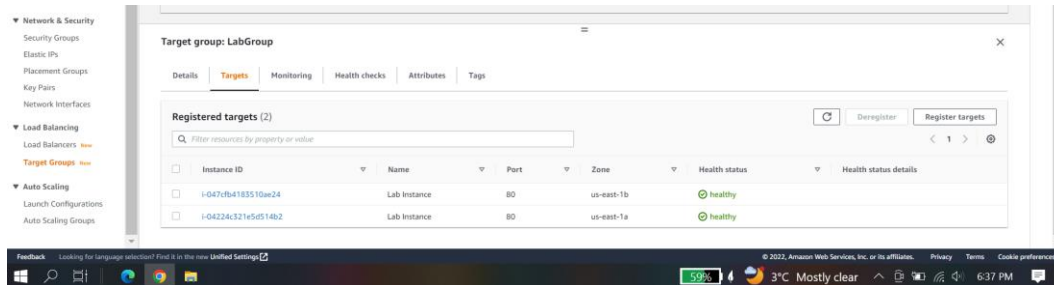
Launch instances

Find instance by attribute or tag (case-sensitive)

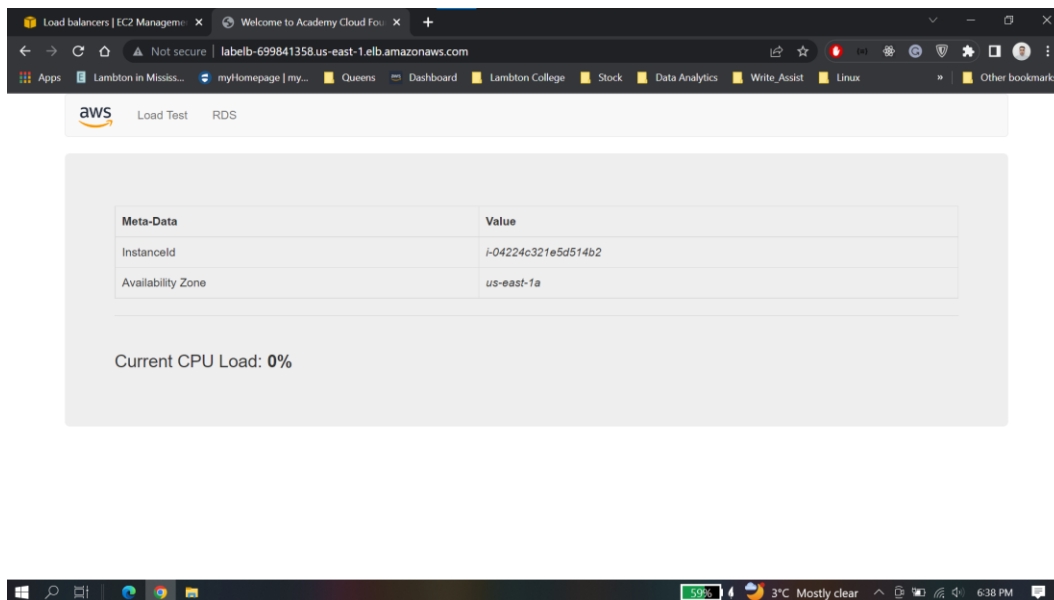
<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP	IF
<input type="checkbox"/>	Bastion Host	i-047fbc8a3985c2c6	Running	t2.micro	2/2 checks passed	No alarms	us-east-1a	-	54.82.76.181	-	-
<input type="checkbox"/>	Web Server 1	i-088763f4dc8dae34a	Running	t2.micro	2/2 checks passed	No alarms	us-east-1a	-	184.73.100.128	-	-
<input type="checkbox"/>	Lab Instance	i-04224c321e5d514b2	Running	t2.micro	Initializing	No alarms	us-east-1a	-	-	-	-
<input type="checkbox"/>	Lab Instance	i-047cfb4183510ae24	Running	t2.micro	Initializing	No alarms	us-east-1b	-	-	-	-

Task 4: Verify that Load Balancing is Working

Healthy indicates that an instance has passed the Load Balancer's health check. This means that the Load Balancer will send traffic to the instance



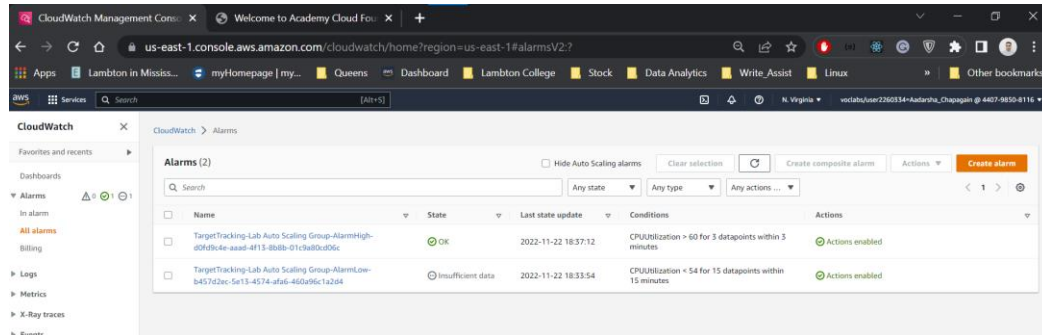
Test Load balancer



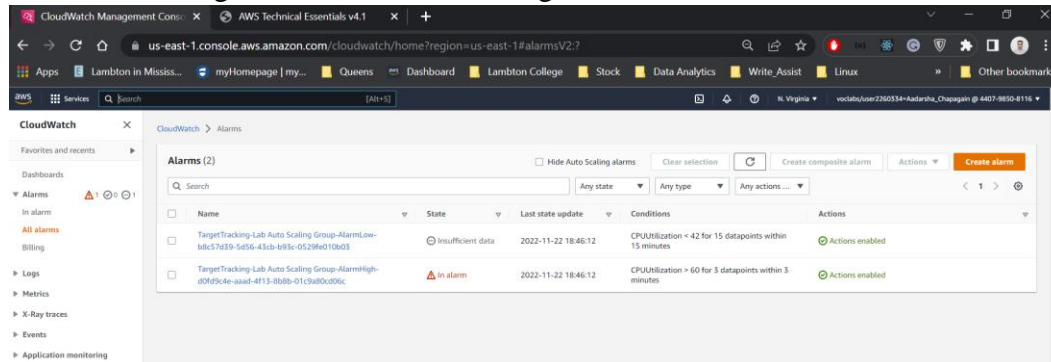
This indicates that the Load Balancer received the request, sent it to one of the EC2 instances, then passed back the result.

Task 5: Test Auto Scaling

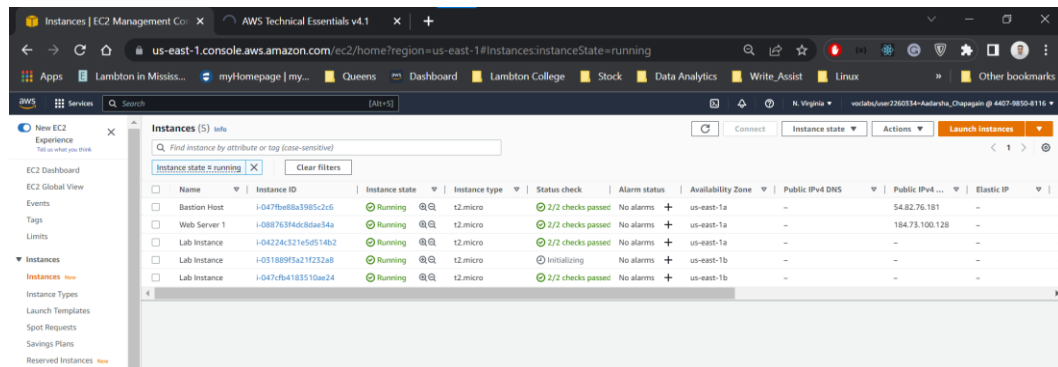
Two alarms under cloudwatch



After Performing load test Here Alarm high is in “**In Alarm**” state



Number of lab instances has increased as well from 2 to 3 instances to handle the load



Lab 6 Conclusion

- Created an Amazon Machine Image (AMI) from a running instance.
- Created a load balancer.
- Created a launch configuration and an Auto Scaling group.
- Automatically scale new instances within a private subnet
- Created Amazon CloudWatch alarms and monitor performance of your infrastructure.